

# Lecture: K-Means Clustering

Rodrigo Canaan

[rodrigocanaan.com](http://rodrigocanaan.com)

[rocanaan@gmail.com](mailto:rocanaan@gmail.com)

 @rocanaan

# Slides and Code

<https://github.com/rocanaan/k-means>



# Types of Machine Learning

- **Supervised Learning**

Learning to predict values or classify objects based on labeled data

- **Unsupervised Learning**

Learning patterns based on unlabeled data

- **Reinforcement Learning**

Learning to act intelligently based on interaction with an environment

# Types of Machine Learning

- **Supervised Learning**

Learning to predict values or classify objects based on labeled data

- **Unsupervised Learning**

Learning patterns based on unlabeled data

- **Reinforcement Learning**

Learning to act intelligently based on interaction with an environment

# Why learn from unlabeled data?



- It's *everywhere*!
- Human labeling is **expensive**!

# **What can we do with unlabeled data?**

- **Gain insight about the data**
- **Compress and/or visualize the data**
- **Generate examples that look like the data**
- **Label data for downstream tasks**

# Examples in various domains

- **Visual Processing**

Group similar images even without labels

- **Recommender Systems**

Group similar products and/or users together

- **Games**

Build models of different player styles

- **Science**

Create a taxonomy of phenomena (e.g. stars) based on their observed properties (magnitude, spectrum, distance...)

# Unsupervised Learning Tasks

- **Clustering**

Partitioning the data into “clusters” based on a measure of similarity

- **Dimensionality reduction**

Transforming the data from a high dimension to a low dimension

- **Generative models**

Learning a “model” that can be used to produce new examples that are similar to the data



# Unsupervised Learning Tasks

- **Clustering**

Partitioning the data into “clusters” based on a measure of similarity

- **Dimensionality reduction**


Transforming the data from a high dimension to a low dimension

- **Generative models**

Learning a “model” that can be used to produce new examples that are similar to the data


# Clustering - Visual Example

Before clustering




# Clustering - Visual Example

Before clustering



After k-means Clustering (k=3)



# Clustering - Distances

- **What does it mean for two data points to be “similar”?**

Usually, it means to have a small “distance” according to some metric


# Clustering - Distances

- **What does it mean for two data points to be “similar”?**

Usually, it means to have a small “distance” according to some metric

**Examples** (in 2-dimensions)

## Manhattan Distance



$$D_M = |x_2 - x_1| + |y_2 - y_1|$$


# Clustering - Distances

- **What does it mean for two data points to be “similar”?**

Usually, it means to have a small “distance” according to some metric


**Examples** (in 2-dimensions)

## Manhattan Distance



$$D_M = |x_2 - x_1| + |y_2 - y_1|$$

## Euclidean Distance



$$D_E = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$$

# Clustering Objectives

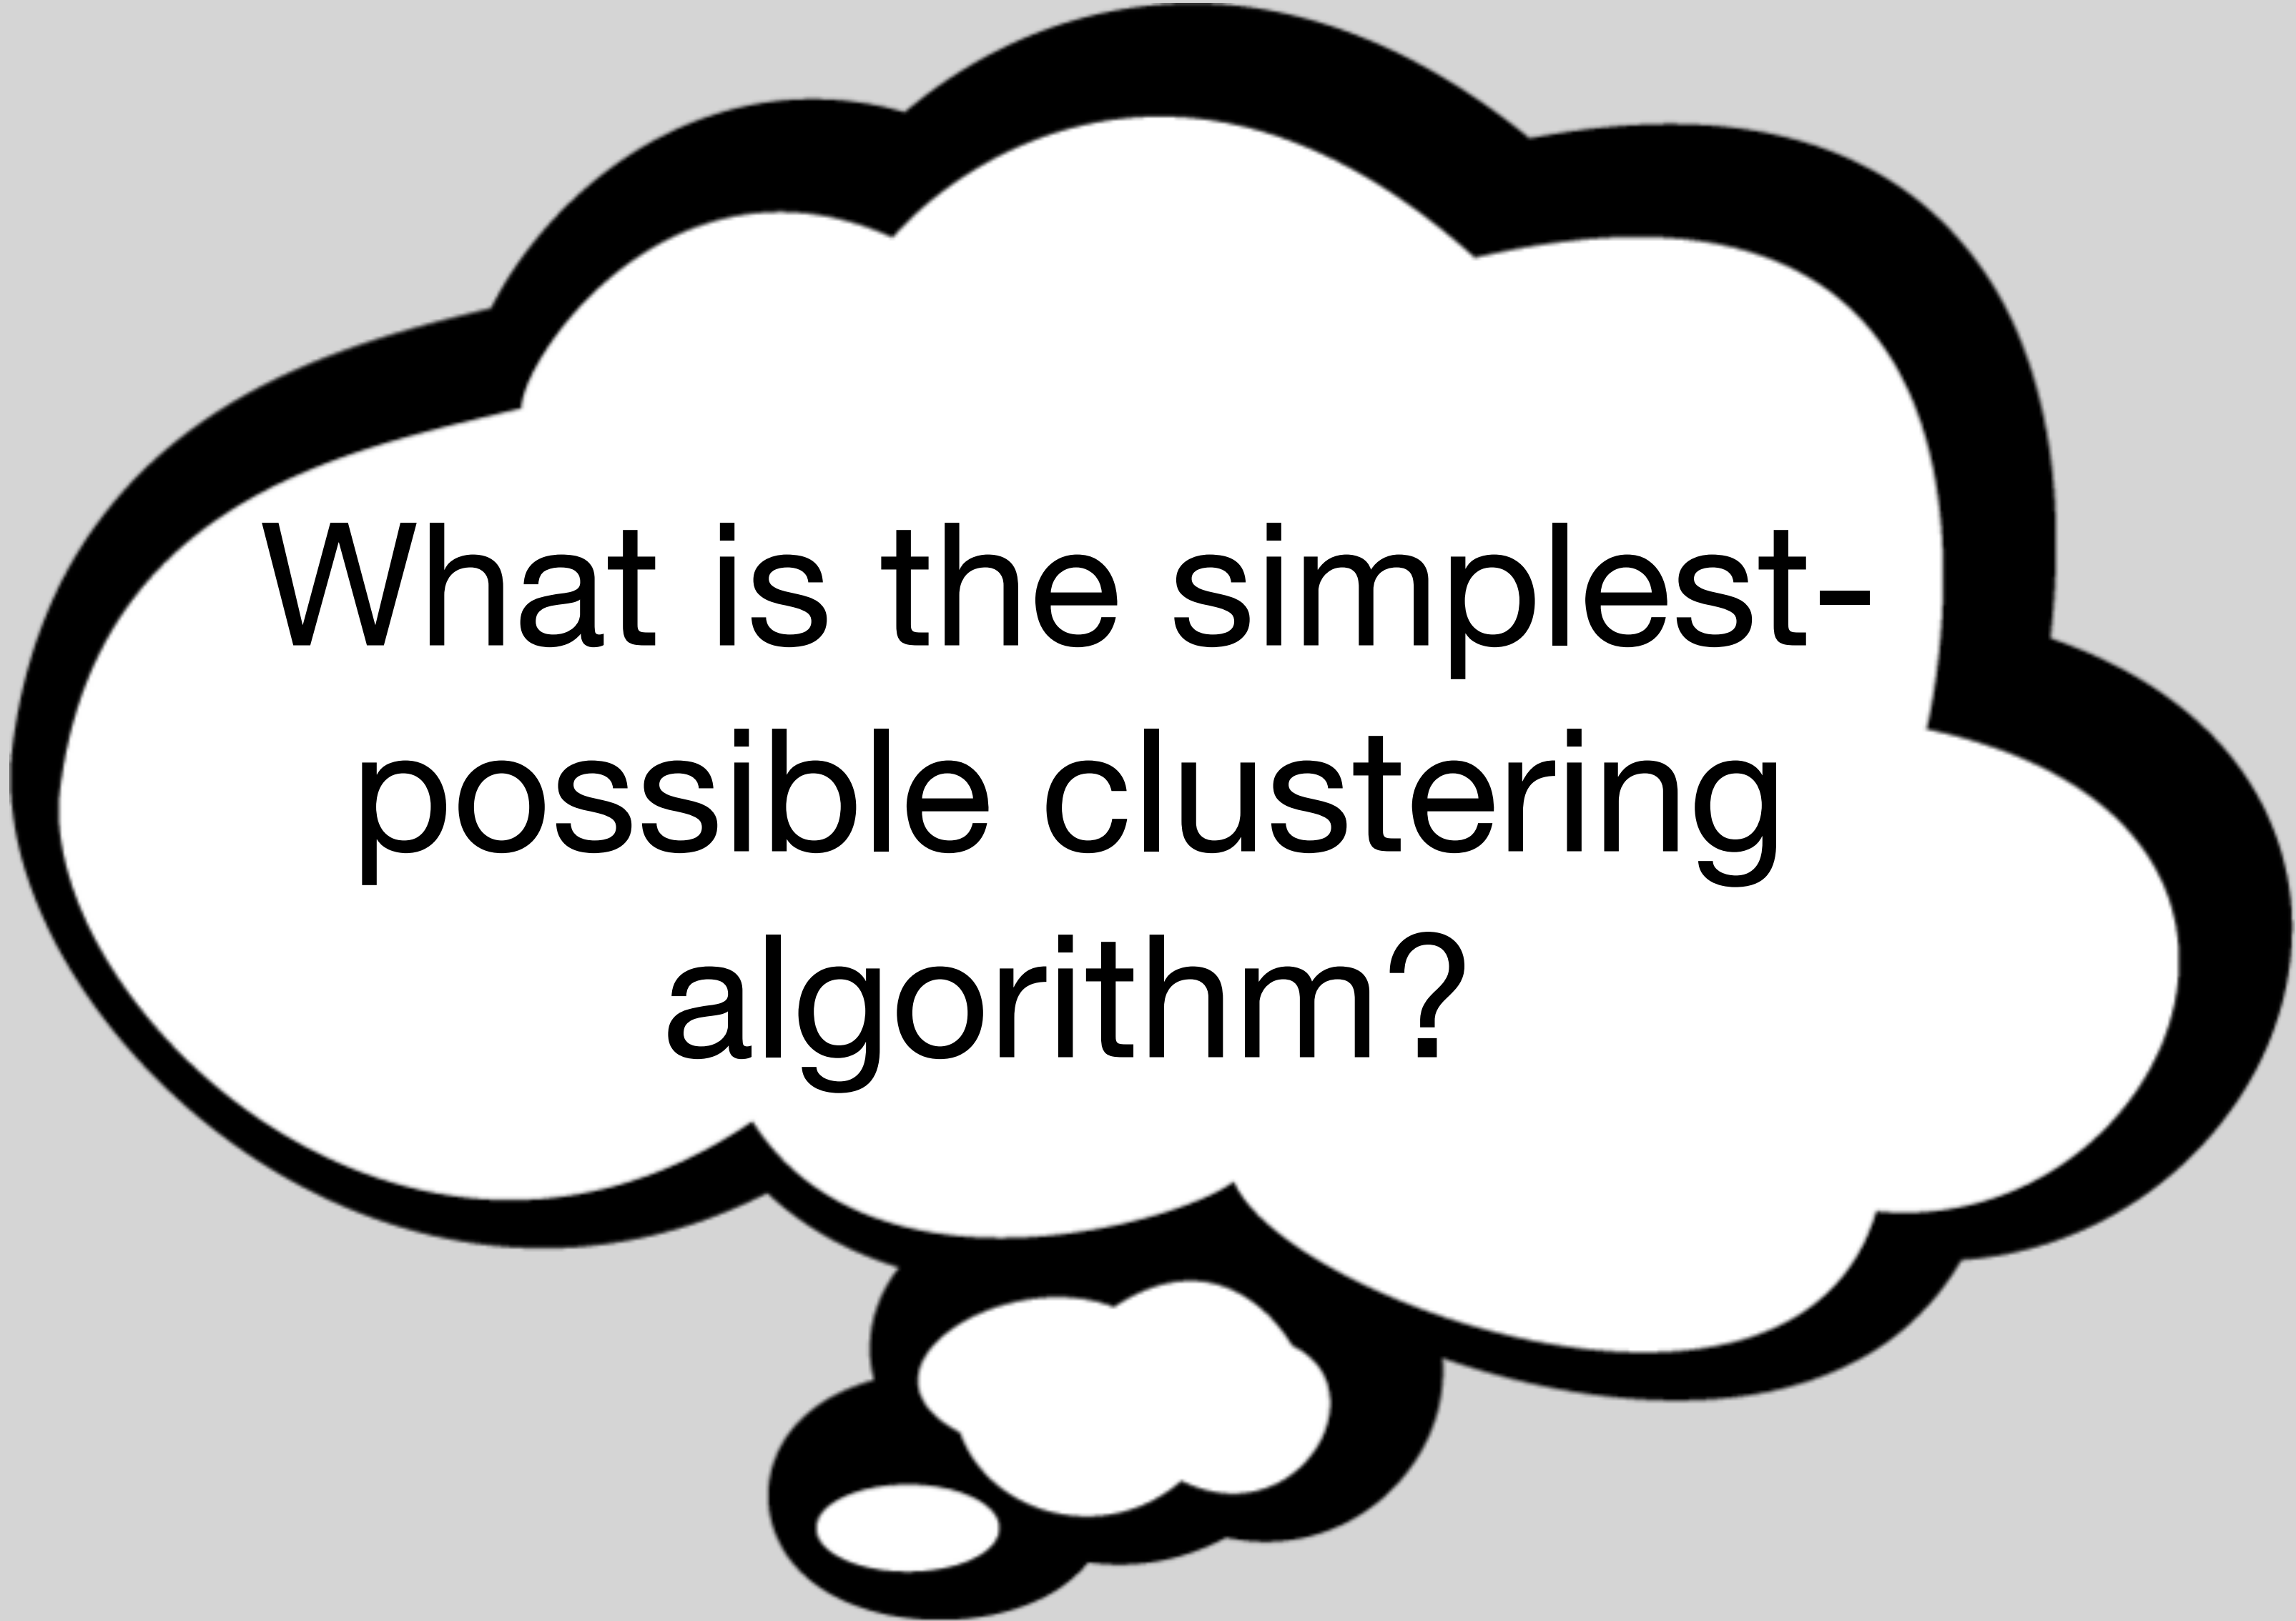
## What does it mean for a clustering to be “good”?

Intuitively:

- Elements in the same cluster are similar (“close”)
- Elements in different clusters are dissimilar (“far”)

### Examples:

- Within-Cluster Sum of Squares (WCSS)
- Average Distance from Centroid
- Maximum Distance from Centroid



What is the simplest-  
possible clustering  
algorithm?



# **Clustering by exhaustive enumeration**

- 1. Enumerate all possible clusterings**
- 2. Evaluate clustering objective for each clustering**
- 3. Return the best clustering**

# Clustering by exhaustive enumeration

1. Enumerate all  $k^n$  possible clusterings
2. Evaluate clustering objective for each clustering
3. Return the best clustering

# Clustering by exhaustive enumeration

1. Enumerate all  $k^n$  possible clusterings
2. Evaluate clustering objective for each clustering
3. Return the best clustering



**Very expensive!**

Need **heuristic** for approximate (but faster) solution!

# Clustering by exhaustive enumeration

1. Enumerate all  $k^n$  possible clusterings
2. Evaluate clustering objective for each clustering
3. Return the best clustering




**Very expensive!**


Need **heuristic** for approximate (but faster) solution!

**Example:** assign each point to cluster with nearest centroid

# Centroid Calculation




# Centroid Calculation



$$x\_mean = (3+4+5)/3 = 4$$

$$y\_mean = (5+7+3)/3 = 5$$

# Centroid Calculation



$$x\_mean = (3+4+5)/3 = 4$$

$$y\_mean = (5+7+3)/3 = 5$$

$$\text{Centroid} = (4,5)$$

# Outline of K-Means Clustering Algorithm

- 1. Generate  $k$  initial centroids**
- 2. Assign points based on heuristic**  
Each point is assigned to cluster with nearest centroid by Euclidean distance
- 3. Recalculate centroids**
- 4. Repeat (2, 3) until no re-assignments**  
Alternatively, up to some maximum number of epochs



# (Naive) K-Means Pseudo-Code

```
1 # Pseudo-code of K-Means clustering algorithm
2 # Assumes a Partition class with methods to maintain centroids and labels
3 # P and P_new are instances of this Partition class
4 Function k-means (data,k)
5     P <- initialize_partition(data,k)
6     stop <- False
7     while not stop # runs until no new assignments
8         # optionally, until some max_iterations is reached
9         P_new <- empty_partition(k)
10        for d in data # Check distance from d to k centroids, assign to closest
11            new_label <- P.get_closest_centroid_label(d)
12            P_new.add_element(d,new_label)
13        Endfor
14        if P_new = P # If nothing changed, stop
15            stop = True
16        Endif
17        P <- P_new
18        P.compute_centroids() # Re-compute centroids based on new labels
19    Endwhile
20
21    return P
22 Endfunction
23
```

# (Naive) K-Means Pseudo-Code

```
1 # Pseudo-code of K-Means clustering algorithm
2 # Assumes a Partition class with methods to maintain centroids and labels
3 # P and P_new are instances of this Partition class
4 function k-means (data,k)
5     P <- initialize_partition(data,k)
6     stop <- False
7     while not stop # runs until no new assignments
8         # optionally, until some max_iterations is reached
9         P_new <- empty_partition(k)
10        for d in data # Check distance from d to k centroids, assign to closest
11            new_label <- P.get_closest_centroid_label(d)
12            P_new.add_element(d,new_label)
13        Endfor
14        if P_new = P # If nothing changed, stop
15            stop = True
16        Endif
17        P <- P_new
18        P.compute_centroids() # Re-compute centroids based on new labels
19    Endwhile
20
21    return P
22 Endfunction
23
```

# K-Means Initialization

## How to initialize clusters?

- **Forgy method:** Choose  $k$  initial centroids randomly (from the data), assign other points according to distance to centroids
- **Random Partition:** Assign each datapoint to a random cluster label, then compute centroids
- Other options: see comparative study by ([Celebi et al., 2013](#))


# K-Means Initialization Pseudo-Code

```
3 # Initial assignment of labels can be done in the main algorithm itself!
4 Function initialize_partitions(data,k)
5     P <- empty_partition(k)
6     P.centroids <- sample_without_replacement(data,k) # get k points from data
7 Endfunction
8
9 #OR
10
11 # Random Partition initialization
12 # Assign each point to a random label, then compute centroids
13 Function initialize_partitions(data,k)
14     P <- empty_partition(k)
15     for d in data
16         label <- random_uniform(k)
17         P.add_element(d,label)
18     Endfor
19     P.compute_centroids()
20     return P
21 Endfunction
22
23
24
25
```




# Demo

Forgy Method, k=3



Random Partition, k=3




# **Advantages and Disadvantages**

# **(Naive) K-Means Disadvantages**

- **Struggles with some data shapes, sizes and outliers**

# (Naive) K-Means Disadvantages


- **Struggles with some data shapes, sizes and outliers**





# (Naive) K-Means Disadvantages


- **Struggles with some data shapes, sizes and outliers**




# (Naive) K-Means Disadvantages

- **Struggles with some data shapes, sizes and outliers**


Non-Convex Clusters



Partition with k=2




Clusters of different sizes




# (Naive) K-Means Disadvantages

- **Struggles with some data shapes, sizes and outliers**


Non-Convex Clusters




Partition with k=2



Clusters of different sizes



Partition with k=3



# **(Naive) K-Means Disadvantages**

- **Struggles with some data shapes, sizes and outliers**
- **Struggles with categorical data**
- **Guaranteed convergence only to local (not global optimum)**
- **Need to specify  $k$  in advance**
- **“Curse of dimensionality”**
- **Worst-case time complexity?**

# Time Complexity

$O(e)$   
Outer-loop  
epochs  
 $O(n)$   
Inner-loop  
iterations

```
1 # Pseudo-code of K-Means clustering algorithm
2 # Assumes a Partition class with methods to maintain centroids and labels
3 # P and P_new are instances of this Partition class
4 Function k-means (data,k)
5     P <- initialize_partition(data,k)
6     stop <- False
7     while not stop # runs until no new assignments
8                     # optionally, until some max_iterations is reached
9         P_new <- empty_partition(k)
10        for d in data # Check distance from d to k centroids, assign to closest
11            new_label <- P.get_closest_centroid_label(d)
12            P_new.add_element(d,new_label)
13        Endfor
14        if P_new = P # If nothing changed, stop
15            stop = True
16        Endif
17        P <- P_new
18        P.compute_centroids() # Re-compute centroids based on new labels
19    Endwhile
20
21    return P
22 Endfunction
23
```

$O(k)$



# Time Complexity = $O(nke)$

$O(e)$   
Outer-loop  
epochs  
 $O(n)$   
Inner-loop  
iterations

```
1 # Pseudo-code of K-Means clustering algorithm
2 # Assumes a Partition class with methods to maintain centroids and labels
3 # P and P_new are instances of this Partition class
4 Function k-means (data,k)
5     P <- initialize_partition(data,k)
6     stop <- False
7     while not stop # runs until no new assignments
8         # optionally, until some max_iterations is reached
9         P_new <- empty_partition(k)
10        for d in data # Check distance from d to k centroids, assign to closest
11            new_label <- P.get_closest_centroid_label(d)
12            P_new.add_element(d,new_label)
13        Endfor
14        if P_new = P # If nothing changed, stop
15            stop = True
16        Endif
17        P <- P_new
18        P.compute_centroids() # Re-compute centroids based on new labels
19    Endwhile
20
21    return P
22 Endfunction
23
```

$O(k)$

Note: number of epochs  $e$  is hard to estimate, can be big in worst case

# **(Naive) K-Means Disadvantages**

- **Struggles with some cluster shapes, sizes and outliers**
- **Struggles with categorical data**
- **Guaranteed convergence only to local (not global optimum)**
- **Need to specify  $k$  in advance**
- **“Curse of dimensionality”**
- **High worst-case time complexity**


# K-Means Advantages

- **Simple to implement**
- **Good performance in many practical scenarios**
- **Adaptations and combinations can handle outliers, different shapes and sizes, higher dimensions, categorical data...**  
Examples: k-medians, k-modes, k-medoids, hierarchical clusterings, kernel methods, dimensionality reduction...
- **Performance can be improved by non-naive implementations**  
Example: using k-d trees to select initial centroids



# Thank you!



[www.rodrirocanaan.com](http://www.rodrirocanaan.com)  
[rocanaan@gmail.com](mailto:rocanaan@gmail.com)  
 @rocanaan

Scholar Page

